

# Creating a highly available and scalable Webservice with OCI

## Important to know for this meetup

login to OCI console:

login url: <https://console.eu-frankfurt-1.oraclecloud.com>

tenant : oci\_core\_emea\_sc\_dolder

user: meetociXX pw: Code4fun.meetoci, XX=01-24

login to wserver vm:

ssh [opc@publicIP](#), authentication by ssh key, get publicIP from OCI console

**NOTICE: we are all operating with the same passwords and keys, so please stay with the user 'meetociXX' that was assigned to you.**

**AD Distribution list for server placement by user, in case we are more than 16 participants**

User	AD1	AD2	AD3
meetoci01	X	X	
meetoci02	X	X	
meetoci03	X	X	
meetoci04	X	X	
meetoci05	X	X	
meetoci06	X	X	
meetoci07	X	X	
meetoci08	X	X	
meetoci09	X		X
meetoci10	X		X
meetoci11	X		X
meetoci12	X		X
meetoci13	X		X
meetoci14	X		X
meetoci15	X		X
meetoci16	X		X
meetoci17		X	X
meetoci18		X	X
meetoci19		X	X
meetoci20		X	X

meetoci21	X	X
meetoci22	X	X
meetoci23	X	X
meetoci24	X	X

## OCI Console LAB:

- Make sure you're working in the us-ashburn-1 region
- Create VCN (simple, no automatic default provisioning, name: **meetociXX-VCN**)
- Create subnet in VCN (either in AD-1 or AD-2, according to instructions), name: **meetociXX-SN1**, use default route table, default security list, default dhcp
- Create internet gateway for this VCN, name: **meetociXX-IGW**
- Create route to IGW in default route
- Create compute instance (shape vm1.2) in AD-1, OL 7.5 latest, name: meetociXX-WS1, configure supplied ssh-keys for subsequent access
- When ready, login to newly created vm 'opc@publicIP' using a key based login
  - Using yum, install 'httpd', create firewall rules, create simple index.html

```
$ sudo yum -y install httpd
```

```
$ sudo systemctl enable httpd
```

```
$ sudo firewall-cmd --permanent --add-key=80/tcp
```

```
$ sudo firewall-cmd --reload
```

```
$ sudo vi /etc/selinux/config (→ SELINUX=disabled)
```

```
$ sudo echo "some text" >/var/www/html/index.html
```

```
c$ sudo vi /var/www/cgi-bin/info.cgi → copy/paste supplied cgi example into file
```

```
$ sudo chmod +x /var/www/cgi-bin/info.cgi → make .cgi executable
```

- Use OCI console to reboot server (needed to disable selinux)
- Test if webserver reachable and working by issuing from browser:
  - publicIP of server → should return "some text" from index.html
  - publicIP/cgi-bin/info.cgi → should return results of cgi script
  - HINT: if it doesn't work, checking route tables (route to IGW?) and security lists (is port 80 traffic into server allowed from the internet)
- If webpages display successfully from the internet, then repeat all steps from above to create a second server in another domain (which really means that they're geographically separated):
  - Basically you create a second subnet in domain AD-2/3, a second server using the AD-2/3 subnet. IGW, security lists, dhcp options we retain, they support region scope.
  - Then do all configuration work until you successfully can display the webpages using the 2<sup>nd</sup> server's publicIP address.

- Create a new security list which will be used along two additional subnets in which the load balancer, name: **meetociXX-LBSEC**, delete all rules such that this list is empty. It will be autofilled by the LB creation process.
- Create another 2 subnets in the same AD's where the server are deployed. Those will 'host' the load balancer. Attach the newly defined security list, default route and dhcp
- Create the loadbalancer, name: **meetociXX-LB** using the newly created subnets, ensure automatic security list setting
- Create backend set (using port 80, protocol http and url '/' for healthcheck) ,then add the OCID's of the two webserver compute instance, *tip: duplicate console tab so you can point one to the servers to use the copy OCID function*, name: **meetociXX-BS**
- Create the listener, http. Port 80, using the newly created 'backendset', name: **meetociXX-LS**
- When all objects ready, note LB publicIP and try in browser 'publicIP' and 'publicIP/cgi-bin/info.cgi' → refreshing should change between the two webserver as the loadbalancer is set to 'round-robin'
  - HINT: if it doesn't work, check security list rules, is INGRESS traffic to the listener authorized

Note: before starting the 'terraform' labs, make sure the interactively created objects are deleted.  
Normally the sequence:

- Terminate loadbalancer
- Terminate the 2 webserver compute instances
- Terminate the VCN

Should completely remove the objects you created.

## Terraform LAB:

- Login to the terraform server (located in OCI) using: [meetociXX@130.61.86.247](mailto:meetociXX@130.61.86.247), pw: Code4fun.meetoci
- Besides 'terraform' your user has also configured the 'OCI CLI' interface, a quick test to check if it works for you is to query all OCI regions that are available by issuing:
  - `oci iam region list`
- to enable 'terraform' operation, some shell variables need to be defined, activate them by issuing in your login directory:
  - `source tf-env-vars`
- change to the tftest directory and issue the following commands to check proper terraform operation:
  - `terraform init`
  - `terraform plan`
  - `terraform apply` → should return the AD-1 identifier of the actual region (no objects created)
- Lab1:

- -quick tf real create functional test creating a VCN
- Copy the .tbe file to the same name without .tbe (thus keeping the template if you messup somehow)
- Apply changes as outlined below
- Lab2:
  - Create a complete loadbalancer setup with terraform.
  - Same copy operation as before to get a .tf file
  - Do all changes as outlined below
  - Create/Inspect/Change/Destroy the loadbalancer using terraform
- Lab3:
  - The lab2 terraform script does not deploy ssh keys to the webserver hosts, so no way you could log into those.
  - Try to find out (online documentation) what needs to be changed to deploy ssh keys and modify the lab2 terraform .tf accordingly
  - If you can't find the solution, the Lab3 folder contains a .tbe which includes the needed changes (a diff of Lab2 and Lab3 .tbe's will reveal)
- Lab4:
  - Further modify the lab2 script to now scale out the webserver infra by adding a third node in the AD-3 which becomes part of the loadbalancer configuration
  - If you can't find the solution, the Lab4 folder contains a .tbe which includes the needed changes (a diff of Lab3 and Lab4 .tbe's will reveal)

### Changes to apply to .tbe files to personalize for your user:

replace:

XX = 2-digit user#

YY= user# \* 10, no leading zeros

D#1 = 0 (index into AD's, returns definition for AD-1)

D#2 = 1 (index into AD's, returns definition for AD-2)

D#3 = 2 (index into AD's, returns definition for AD-3)

### Additional things to try:

- Delete one of the servers using 'terraform delete target='resoruce'
- Use 'terraform refresh' to view new status
- Check in browser that webpage still works but without deleted server
- Check with 'terraform plan' if a subsequent 'apply' would re-add deleted objects
- Run re-add